

Dynamic Dataflow

Risto Vuorio

Abstract—Dynamic dataflow is a model of computation that is well suited for the construction of parallel programs. In dynamic dataflow the program logic is divided into actors that can execute in any order depending only on the data availability. In most dynamic dataflow implementations the actors are stateless, which results in a model of computation that resembles functional programming in many respects.

As the need for high performance parallel computing is growing, dynamic dataflow based frameworks such as TensorFlow have been introduced to make programming parallel programs easier. In this paper an introduction to the dynamic dataflow model of computation is given, the motivation for adoption of the model for parallel computing is explained and example implementations of the model are explored.

I. INTRODUCTION

A program following the dataflow model of computation (MoC) consists of a directed graph with data flowing between the nodes. The data is split into tokens that are passed between the nodes. The execution of the nodes is asynchronous. Due to the asynchronous execution the tokens have to be buffered between the nodes. The dataflow MoC allows for unbounded execution of the model, which means the dataflow program may execute for a very long time. The possibility of unbounded execution leads to the problems the different dataflow MoCs try to solve. A dataflow program capable of unbounded execution 1. may have unbounded buffer growth 2. if there are cycles, the execution may result in a deadlock where there are not enough tokens to advance the execution. A good introduction on the dataflow MoCs is given in [1].

One popular approach to solving these problems is the synchronous dataflow (SDF). In SDF the number of tokens produced and consumed by each actor is fixed. The SDF MoC guarantees bounded buffers and deadlock-free execution but it is very constrained. For expressing more complicated programs, models with more design freedom are needed. [1]

There exists a variety of Dataflow MoCs that extend the concept of synchronous dataflow such as the parameterized and interfaced synchronous dataflow (PiSDF) [2], which extends SDF expressive power by defining parameters and interfaces. The resulting PiSDF model can be expressed as a SDF. We will not look at these extensions of SDF but at the more generic Dataflow MoCs categorized under dynamic dataflow. Dynamic dataflow does not refer to a single MoC but is rather an umbrella term under which many MoCs fall.

II. DYNAMIC DATAFLOW

In Dynamic Dataflow (DDF) the number of tokens consumed or produced by an actor in a single firing is not constrained. An actor can produce and consume different number of tokens in every firing. This freedom improves the

expression power of the model but makes the analysis more difficult. The difficulty is underlined by the fact that for the most general class of dataflow models that bounded buffers and deadlocks are not decidable as proved by Buck [3]. It is possible to define dynamic dataflow models that are decidable by introducing limitations on the types of actors and graph patterns that can be used in the model [4], [5]. However these limitations are most often traded for the increased expressive power of the models without such limitations [4].

Bhattacharyya et al. [4] describe many examples of DDF MoCs. One of these examples is the CAL Actor Language (CAL). CAL is used for example by the MPEG Reconfigurable Video Coding library. Another example use of dynamic dataflow is in the TensorFlow (TF) machine learning library by Google [6]. TF programs are structured as DDF graphs. The computations in the nodes can be distributed to heterogeneous computing devices such as CPUs and GPUs. TF provides control flow operators that can be added to the graph to support conditional execution of parts of the graph and loops.

Dynamic dataflow is a useful model of computation for handling streaming data. Signal processing applications typically deal with streaming data in the form of audio, video or other kinds of streams. Streaming data can also be other types of numeric data, as in the TensorFlow framework, or even textual data as is often processed in web analytics. The successive and independent dataflow actors provide logical division of computation for stream processing. The stream processing context is encountered in the practical use of DDF MoCs [6], [7] and in the research where most of the studies are focused on the performance of either video or audio stream applications [4], [8], [9].

Although some of the practical implementations of dynamic dataflow [6], [7] allow stateful actors, most of the dynamic dataflow MoCs presented in [4] do not implement stateful actors to help the analysis of the models. The division of computation to stateless and independent blocks resembles functional programming to some extent. This similarity between functional programming and dataflow models of computation is explored in [10].

A. Common Features of DDF Models

The dynamicity of dynamic dataflow models is introduced through variable number of input tokens consumed and output tokens produced by the actors [4]. In stateless actors the number of inputs consumed and outputs produced depends solely on the inputs. The simplest dynamic actors are the switch and select actors presented in figure 1. The select actor takes two or more data inputs and a control input. The actor selects which input to forward to the output port based on the




Fig. 1. Dynamic dataflow switch and select actors. Visualisation from [1]




Fig. 2. Dynamic dataflow conditional constructed from the switch and select actors. Visualisation from [1]

control input. The switch actor is the counterpart for select, taking only one data input and a control input. The switch actor outputs its input tokens to the output port selected with the control input. [1] In the switch and select actors presented in the figure the number of tokens produced and consumed on each port is always zero or one but it could be more in some other kind of actor.

From simple conditional actors presented in figure 1 dynamic data paths can be constructed. An example of conditionally executed actors is presented in figure 2. In the figure actors C and D are conditionally executed based on the control tokens received from the actor B. The conditionals provide means for implementing loops and optional data paths in the dataflow models. These nontrivial extensions improve the expressive power of the dynamic dataflow models.

B. Scheduling the dynamic dataflow graphs

As was stated earlier the improved expressive power of the dynamic dataflow models results in problems with unbounded buffers and the schedulability of the graphs. Buck [3] defines scheduling in the context of dataflow graphs as consisting of three operations. 1. Assigning actors to processors 2. determining the order of execution of the actors on each processor 3. determining the exact starting time of each actor. The strength of synchronous dataflow models is that these decisions can be made at compile time. The dynamic dataflow models need to implement a run time scheduling system. Especially in signal processing applications the reduced run time overhead and analysability of compile time scheduling is considered a big advantage. Therefore there exists dataflow models of computation that combine compile time and run time scheduling to achieve some of these advantages [3], [4].

The dynamic scheduling of the dataflow graphs depends on the specific dataflow model. The basic principle of dynamic scheduling are that actors may execute after their inputs are available. A naive scheduler could be implemented as a round-robin that checks the state of each actor in the graph and executes any actor that has its inputs ready. However this type of scheduler results in inefficient execution as it does not utilize the structure of the graph to make scheduling decisions. More sophisticated scheduler implementations are provided in the different DDF implementations. For example in CAL the graph is first partitioned in to subgraphs that can be statically scheduled. In some cases where there are a limited amount of alternative data paths the scheduler may generate multiple static schedules involving some of the same actors. The dynamic scheduling decisions are then made between these partitions. [7]

C. Example problems solved using DDF

The research on dataflow models of computation have been studied extensively for decades and especially synchronous dataflow has been used in practice as well. It is no wonder then that the research community has adopted a set of characteristic examples of the model of computation. The MPEG decoder seems to be the most commonly studied workload for any new dynamic dataflow model. It is implemented in CAL [7] and SADF [4] to name two examples. MPEG decoder is a well defined but highly complex stream processing application with practical relevance so it lends itself well to benchmarking purposes. The MPEG video format achieves compression of the video data by using motion estimation.

Roughly the video stream is compressed with the following technique. The stream is encoded into reference frames which represent single frames of the video stream and predictive-coded frames which represent change from the previous frames. To decode the stream the difference frames need to be applied on the reference frames. Dynamicity in the dataflow model of the decoder is introduced after the input frame is parsed and the type of the frame is determined. The reference frames and the difference frames are processed differently. Detailed descriptions of partitioning the MPEG4 decoder into dynamic dataflow graph are given in [7], [8].

Other example uses of DDF models include encoding and decoding of MP3 audio format in [4] and computing live statistics on arriving tweets in [11].

D. TensorFlow

TensorFlow is a Google framework for machine learning on heterogeneous computing platforms [6]. The motivation for a variant of dataflow computing derives from two factors. First the need for efficient means of construction of parallel programs is only highlighted on heterogeneous platforms where the computations may take place on different types of computational units such as CPUs and GPUs and on distributed nodes that are physically separated. For practical usability machine learning algorithms need to be able to handle massive datasets. Many machine learning algorithms break




Fig. 3. A partial dataflow graph of a simple neural network. The inputs b , W and x are at the bottom of the graph. The rounded rectangles represent computation nodes and the output C is at the top. The graph is taken from [6]

down into sequential phases of computation that translate into actors naturally. An simple partial example of a machine learning algorithm represented as a dataflow graph is presented in figure 3.

The dataflow MoC in TensorFlow is dynamic. Some of the TensorFlow nodes maintain and update a persistent state. TensorFlow supports loops in the graph, branching of the graph, conditional execution of nodes and control dependencies. Control dependencies are edges between the nodes with no data flowing along but they impose rules on the execution order. If there is a control dependency from node A to node B, B can not execute before A has completed. [6]

The TensorFlow runtime schedules the actors on compute devices it has been configured to use automatically. The open source distribution supports CPUs and GPUs out-of-the-box. The distributed execution on different host machines is not currently supported by the open source version. User has additional control of the asynchronous execution of sub-graphs through the use of queues. The user can define enqueue and dequeue operations in the graph to utilize queueing features of the framework. The queues can be used for example for pre-fetching data from the disk while other part of the graph is using the computational resources. [6]

III. CONCLUSION

In this paper an overview of the dynamic dataflow models of computation was given, the common structures between the models were discussed and some application areas of the models were introduced. A wider look at the DDF MoCs is given in [4].

Dynamic dataflow is not a single strictly defined model of computation. Many different models fall under the classification but they share a set of features. All DDF models are based on dataflow graphs. Dataflow graphs are constructed of actors

connected by edges along which tokens of data flow. While the different DDF MoCs have different rules for which kinds of actors are allowed, how the actors are scheduled etc., they all allow some data dependencies in the graph which results in requirement for dynamic scheduling. This balance between expressive power and analysability of the model is optimized by each model according to its goals.

The demand for intuitive and powerful ways to describe parallel computations is growing. One answer to the demand is the use of dynamic dataflow models of computation. DDF MoCs have been a subject of research for decades but have failed to garner widespread support among practical users. With the introduction of new DDF frameworks such as the ones presented in [6], [11] the DDF models have a chance of making it to the mainstream.

REFERENCES

- [1] E. A. Lee and S. A. Seshia, *Introduction to embedded systems: A cyber-physical systems approach*. Lee & Seshia, 2015.
- [2] K. Desnos, M. Pelcat, J.-F. Nezan, S. S. Bhattacharyya, and S. Aridhi, "Pimm: Parameterized and interfaced dataflow meta-model for mpsoes runtime reconfiguration," in *Embedded Computer Systems: Architectures, Modeling, and Simulation (SAMOS XIII)*, 2013 *International Conference on*. IEEE, 2013, pp. 41–48.
- [3] J. T. Buck, E. Lee *et al.*, "Scheduling dynamic dataflow graphs with bounded memory using the token flow model," in *Acoustics, Speech, and Signal Processing*, 1993. *ICASSP-93.*, 1993 *IEEE International Conference on*, vol. 1. IEEE, 1993, pp. 429–432.
- [4] S. S. Bhattacharyya, E. F. Deprettere, R. Leupers, and J. Takala, *Handbook of signal processing systems*. Springer Science & Business Media, 2013.
- [5] G. Gao, R. Govindarajan, and P. Panangaden, "Well-behaved dataflow programs for dsp computation," in *Acoustics, Speech, and Signal Processing*, 1992. *ICASSP-92.*, 1992 *IEEE International Conference on*, vol. 5. IEEE, 1992, pp. 561–564.
- [6] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat, I. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia, R. Jozefowicz, L. Kaiser, M. Kudlur, J. Levenberg, D. Mané, R. Monga, S. Moore, D. Murray, C. Olah, M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. Tucker, V. Vanhoucke, V. Vasudevan, F. Viégas, O. Vinyals, P. Warden, M. Wattenberg, M. Wicke, Y. Yu, and X. Zheng, "TensorFlow: Large-scale machine learning on heterogeneous systems," 2015, software available from tensorflow.org. [Online]. Available: <http://tensorflow.org/>
- [7] J. Eker and J. Janneck, "Cal language report," Tech. Rep. ERL Technical Memo UCB/ERL, Tech. Rep., 2003.
- [8] G. Roquier, M. Wipliez, M. Raulet, J. W. Janneck, I. D. Miller, and D. B. Parlour, "Automatic software synthesis of dataflow program: An mpeg-4 simple profile decoder case study," in *Signal Processing Systems, 2008. SiPS 2008. IEEE Workshop on*. IEEE, 2008, pp. 281–286.
- [9] J. Ersfolk *et al.*, "Scheduling dynamic dataflow graphs with model checking," 2014.
- [10] H. J. Reekie, "Realtime signal processing—dataflow, visual, and functional programming," 1995.
- [11] D. G. Murray, F. McSherry, R. Isaacs, M. Isard, P. Barham, and M. Abadi, "Naiad: a timely dataflow system," in *Proceedings of the Twenty-Fourth ACM Symposium on Operating Systems Principles*. ACM, 2013, pp. 439–455.